

CS5223

Distributed Systems

Lecture 0: Course Introduction

Instructor: Haifeng YU

Course Overview

- Homepage: go to IVLE and search for CS5223
 - Check course homepage often!
 - General student discussion forum
 - Prerequisite:
 - CS2106 & CS3211
- OR
- Basic course in operating systems, knowledge in parallel/concurrent programming (i.e., locks, monitors, etc), knowledge in programming in Java
 - If you only know how to program in C but not Java, you may also enroll in this course but expect to teach yourself some Java along the way
 - If not sure, come and ask me...

Course Overview

- Main objectives:
 - Learn basic concepts of distributed systems
 - Learn how to design and build distributed systems
 - Learn how to reason about the correctness of distributed systems
- Side objectives (achieved along the way of achieving main objectives):
 - Learn about basic methodologies in how to do research
 - Get into contact with various research-related activities (e.g., literature survey)

Syllabus

- Tentative lesson schedule available on IVLE
- Take a look to see what topics will be covered
- Take a look at the assignments

Is This Course Hard?

- It is not a course that you can pass without spending substantial amount of time – not a leisure course with an “easy-way-out”
 - Substantial efforts and time commitments are expected
- You will need to do **substantial** software building/implementation (i.e., coding) in this course

AND

- You will need to be able to think abstractly about the protocols and the algorithms

both theory and practice are important in distributed systems

- If you are not comfortable with either, re-think before you take this course

Course Format

- Instructor: Haifeng YU (office COM2-04-25)
- Required text book:
 - “Distributed Systems – Principles and Paradigms” by Andrew Tanenbaum and Maarten Van Steen (2nd edition / 2006 – no newer editions available)
- Supplementary reading:
 - “Distributed Systems: Concepts and Design” by George Coulouris, Jean Dollimore, Tim Kindberg (5th edition / 2011)
- No perfect textbook on this subject
- Office hour: Monday 1:30pm—3:30pm every lecturing week
- Read material before each lecture (~3 hours)
 - You won't follow if you don't read beforehand
- Assignments/Homeworks (On average ~5 hours/week)

Class Duration

- Every Monday 6:30pm – we will start on time, so don't be late
- NUS has the following guideline on class duration:
 - “With effect from AY 2011/2012 ... For 2-hour lessons, Faculty Members are requested to keep the break to 5 minutes, and to release students 25 minutes before the end of the second hour. ... We request the co-operation of all Faculty Members to observe the above policy.”
 - Above text taken from Provost's Circular No. 3 of Academic Year 2011/2012 dated 8 September 2011, by NUS Provost.
- Thus
 - We will have a 5-minute break in the middle
 - Class will end on 8:05pm

Grading Policy

- 30 marks: Assignment One (1 September 2014 to 9 October 2014)
 - Design and implementation of a small distributed system
- 30 marks: Assignment Two (13 October 2014 to 10 November 2014)
 - Literature survey on a research subject
- 40 marks: Final exam (Open Book)
- Both assignments are online already
 - Take a look to get better idea about what is expected
 - Assignments have been changed/updated from previous years...so don't expect them to be the same as what your seniors did

Team Forming

- Both Assignment One and Assignment Two are team assignments
 - Each team has at most 2 students
 - You cannot have the same teammate for the two assignments
- Team forming deadline (for both assignments):
11:59PM on Wednesday 27 August 2014
 - Email (in one email) the TA (zhouzl@comp.nus.edu.sg) your name and your teammates' name for each assignment
 - 1 point penalty for each assignment per day after 27 August 2014
 - The penalty serves to prevent some people from waiting until the last minute to form teams

Demo for Assignment One

- 11-12 October 2014 (Saturday and Sunday) will be the demo days for Assignment One
 - We will assign you a time slot
 - I expect everyone to show up during the assigned slot on that day – this should be viewed as a pre-condition for taking this module (similar to the date for final exam...)
 - **Speak up NOW if you have difficulty coming...**
- **Repeat: This is a pre-condition for taking this module!**
 - Will NOT honor any request for doing the demo on a different time/date – this is a commitment
 - No special consideration will be given (and NEVER had been given before in previous years)

Small Amount of Informal Homework

- Does not count toward final grade
- No need to submit – will be discussed in class
- But why you should do it
 - Help you to do better in final exam
 - An important and prompt feedback to you
- Both Assignment One and Assignment Two are due only after half of the semester is gone
 - You may realize too late that you are not following (if you ignore the homework)

Class Participation and Penalty for Non-participation

- The School and also myself encourage class participation
 - But hard to incorporate into assessment (fairness issues)
 - No webcast
- My approach to encourage participation:
 - If you miss any important discussions/announcements I made during lecture – you pay the price yourself
 - There will be important things that are only discussed during lecture, and not via email or IVLE or other venues
- Every student should be committed to attend every lecture
 - This is a requirement for taking this module – do not take this module if you cannot meet this requirement

No Exam-Oriented Training

- You are familiar with:
 - Preparation classes for GRE, Toefl, Gmat, etc.
 - High school classes to prepare you to score better in university entrance exams
 - But this course is very different...
- Final exam is the only paper-pencil exam for this course
 - Preparing for the final exam is not our (most important) goal

Why not train students to get better scores in the final exam?

- Such training is **useless for your future** (after this course)
 - You will likely not take another exam in distributed systems for the rest of your life...
- Such training is **useless for your present** (for this course)
 - Your final grade is largely about your **relative performance**, as compared to your classmates
 - Training may help you to get higher mark, but probably not better grade

On Cheating...

- I hope this slide is never needed...

ZERO TOLERANCE POLICY:

Cheating \Rightarrow Reported to Dept and School

- I keep my promises:
 - I reported multiple plagiarism cases last year
- Example:
 - Copy code/text from other people (e.g., from the Internet or from any other student)

On Experiences from Senior Students

- Some students like to get “pragmatic advices” on how to “game” a module’s policy to get a better grade, from students who took the same module before
- I generally discourage that
 - Sometimes these pragmatic advices do not help student learning
 - Sometimes this causes unfairness among students
- How do I discourage this?
 - I try to change assessments from year to year – so the “pragmatic advices” you get by treating your seniors dinner might not be applicable 😊
- If however, you can get good advices on how to learn more from this module, by all means take those...

Question?



Survey Time

- A 5000 level course is always difficult to teach because of the diversity in background
 - Some may already know a lot about distributed systems (may get bored)
 - Other may not know anything (may get confused)
- The democratic solution – survey
- We need to reschedule our class on 6 October 2014, since it is a public holiday
 - We will do a make-up lecture on 27 Aug or 28 Aug or 29 Aug
 - Make your choice in the survey